

CS523: Deep Learning

ImagePrompt: Discrete Prompt Optimization for Image Generation using Policy Gradient Methods

Group Members:

Reana Naik

Sandesh Bharadwaj

Pruthvi Bharadwaj

Overview

- Motivation & Problem Formulation
- Architecture
- Methods
- Results

Motivation & Problem Formulation

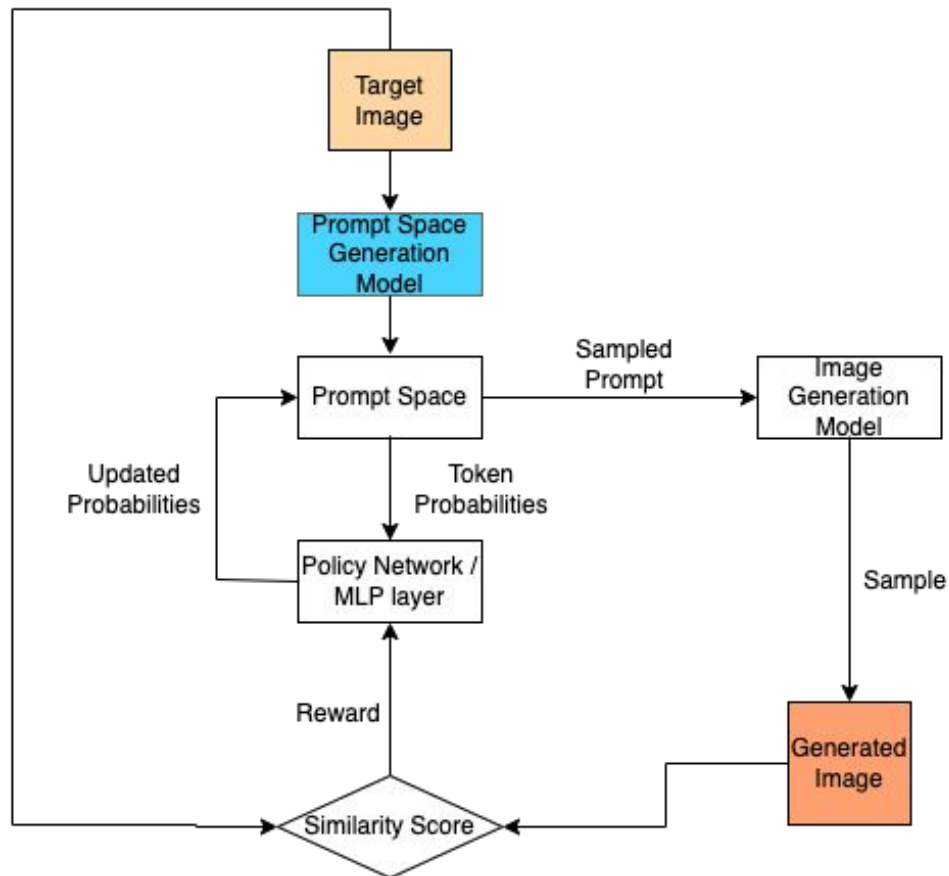
- Forensic artists can be replaced with generative AI models
- Prompt engineering to get the generated image to match the target image
- Prompt Engineering - Generating desired prompts as required for text-to-image and language models.
 - ❖ Time-consuming
 - ❖ Labor-intensive
 - ❖ Size of optimal prompts for images can vary based on image.
- Image captioning - Generating a human-readable text description of a given image.

Objective:


- Generate an optimal prompt from captions generated for an input image using reinforcement learning, such that the prompt generates an image with maximum similarity to the input image.

Methods: *ImagePrompt Architecture*

Architecture of ImagePrompt



Methods: *Generating the Prompt Space*



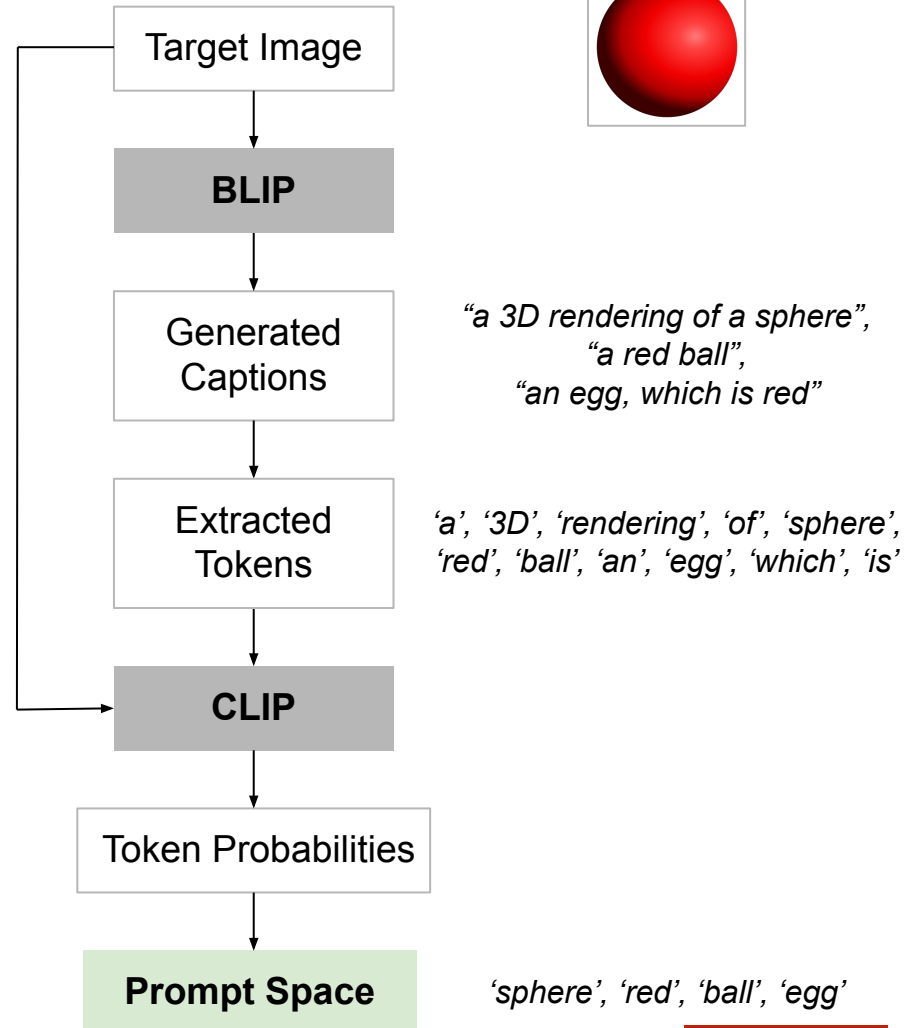
BLIP generates captions for input image.

Extract unique tokens from all generated captions.

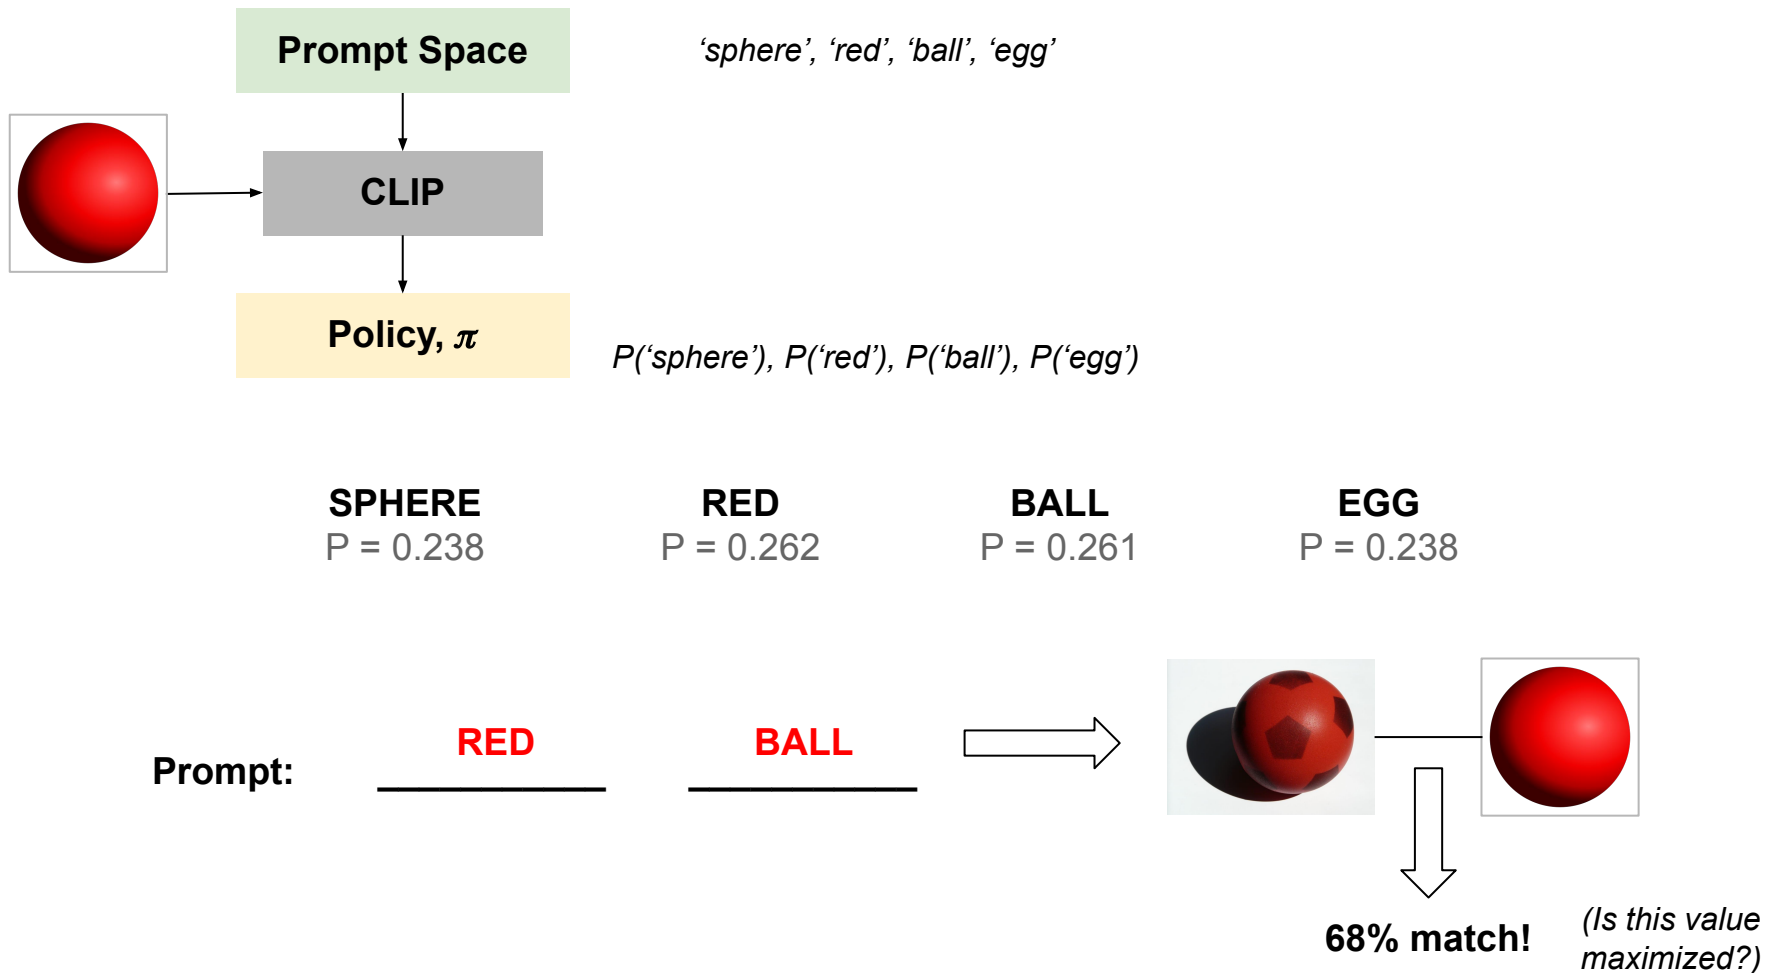
CLIP extracts token and image features for every token and input image.

Calculate similarity between image and token features to get token probabilities

Choose set number of tokens with highest probabilities.



Methods: *Policy Optimization Task*



Methods: *Policy-Gradient Method (REINFORCE)*

REINFORCE is a policy-gradient method that learns the optimal policy directly (on-policy). The parameter being updated are the weights θ of the neural network.

Instead of calculating the reward after every episode, we generate a *trajectory*:

$$\tau = (s_0, a_0, r_1, s_1, a_1, r_2, \dots, a_k, r_{k+1}, s_{k+1}, \dots)$$

A trajectory consists of all the transitional states, actions and rewards in each episode.

We calculate the *return* for a trajectory τ as $R(\tau)$:

$$R(\tau) = G_1 + G_2 + \dots + G_k$$

where G_k is the future return at time step k for a transition episode k

$$(s_k, a_k, r_{k+1})$$

We optimize θ by maximizing the expected return $U(\theta)$ defined as:

$$U(\theta) = \sum_{\tau} P(\tau; \theta) \cdot R(\tau)$$

where $P(\tau; \theta)$ is the probability of each possible trajectory.

Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261


4
EGG
P = 0.238

Prompt: **RED**

Token History:

Reward History:

State History:



Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261

4
EGG
P = 0.238

Prompt: **RED** **SPHERE**

Token History: [2]

Reward History: [0.49]

State History:



Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261


4
EGG
P = 0.238

Prompt: **RED** **SPHERE**

Token History: [2, 1]

Reward History: [0.49, 0.89]

State History:



Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261


4
EGG
P = 0.238

Prompt: EGG _____

Token History: [2, 1]

Reward History: [0.49, 0.89]

State History:



Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261



4
EGG
P = 0.238

Prompt: **EGG** **RED**

Token History: [2, 1, 4]

Reward History: [0.49, 0.89, 0.43,]

State History:



Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261


4
EGG
P = 0.238

Prompt: EGG RED

Token History: [2, 1, 4, 2]

Reward History: [0.49, 0.89, 0.43, 0.91]

State History:



Methods: *Collecting Trajectory Data*

1
SPHERE
P = 0.238

2
RED
P = 0.262

3
BALL
P = 0.261


4
EGG
P = 0.238

Prompt: EGG RED

Token History: [2, 1, 4, 2]

Reward History: [0.49, 0.89, 0.43, 0.91]

State History:



Future Returns: [0.49, 1.38, 0.92, 1.34]

$$R(\tau) = G_1 + G_2 + \dots + G_k$$

$$\tau = (s_0, a_0, r_1, s_1, a_1, r_2, \dots, a_k, r_{k+1}, s_{k+1}, \dots)$$

Methods: *Policy Network Training*

```
model = nn.Linear(512, 4)
optimizer = torch.optim.Adam(net_parameters(), lr=lr)

for _ in range(n_iterations):

    observations, actions, future_returns = ...

    logits = net(observations)
    policy_distributions =
    torch.distributions.Categorical(logits=logits)
    log_probs = policy_distributions.log_probs(actions)

    mean = future_returns.mean()
    std = future_returns.std().clamp_min(1e-12)
    normalized_future_returns = (future_returns -
    mean)/std

    loss = -(log_probs * normalized_returns).mean()

    net.zero_grad()
    loss.backward()
    optimizer.step()
```

*Initialize the policy network
Choose optimizer as Adam*

*Collect trajectory data for
multiple episodes*

*Compute the policy distribution
for each observation and the
log probabilities of each action
as determined by the policy
distribution.*

*Normalize future returns to
reduce variance*

Loss minimization

*Zero the gradients,
propagate the error
backwards and update θ*

Results: *Prompt Generation with Target Image Only* (Policy Gradient Method)

```
# Hyperparameters
```

```
n_captions = 5, n_tokens = 4
```

```
input_size = 512, output_size = 4, lr = 0.001
```

```
max_episodes = 10, n_iterations = 40
```

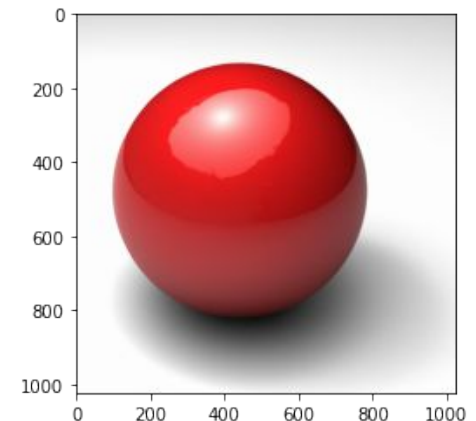
```
prompt_length = 2
```

```
Prompt Space: ['sphere', 'egg', 'ball', 'red']
```

```
Prompt Space Probabilities: [0.23805307 0.23905936 0.2610057  
0.26188195]
```

```
Generated Prompt: red sphere
```

```
Similarity Score: 0.93471456
```



Results: *Prompt Completion with Target Image and Partial Prompt (Policy Gradient Method)*

```
# Hyperparameters
```

```
n_captions = 5, n_tokens = 4
```

```
input_size = 512, output_size = 4, lr = 0.001
```

```
max_episodes = 20, n_iterations = 20
```

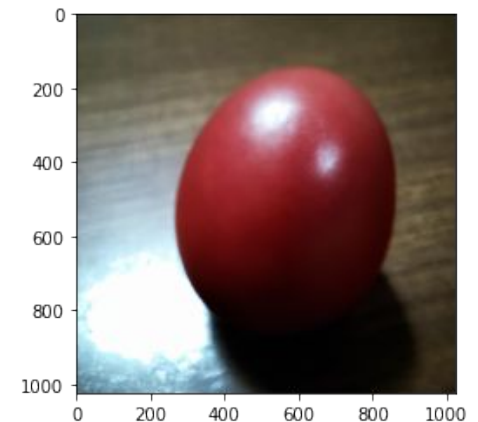
```
prompt_length = 2, init_prompt = 'a red'
```

```
Prompt Space: ['sphere', 'egg', 'ball', 'red']
```

```
Prompt Space Probabilities: [0.23805307 0.23905936 0.2610057  
0.26188195]
```

```
Generated Prompt: a red egg
```

```
Similarity Score: 0.8069409
```



Results: *Prompt Generation with Initial CLIP Probabilities*

```
# Hyperparameters
```

```
n_captions = 5, n_tokens = 4
```

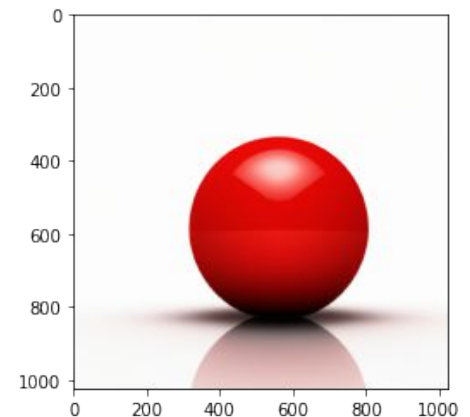
```
prompt_length = 2
```

```
Prompt Space: ['sphere', 'egg', 'ball', 'red']
```

```
Prompt Space Probabilities: [0.23805307 0.23905936 0.2610057  
0.26188195]
```

```
Generated Prompt: red ball
```

```
Similarity Score: 0.9255184
```



Results: *Prompt Generation with Target Image Only* (Policy Gradient Method)

```
# Hyperparameters
```

```
n_captions = 8, n_tokens = 7, max_episodes = 10, n_iterations = 20
```

```
prompt_length = 4
```

```
Prompt Space:
```

```
['the', 'peaks', 'mountain', 'spain', 'horse', 'mountains', 'horses']
```

```
Prompt Space Probabilities:
```

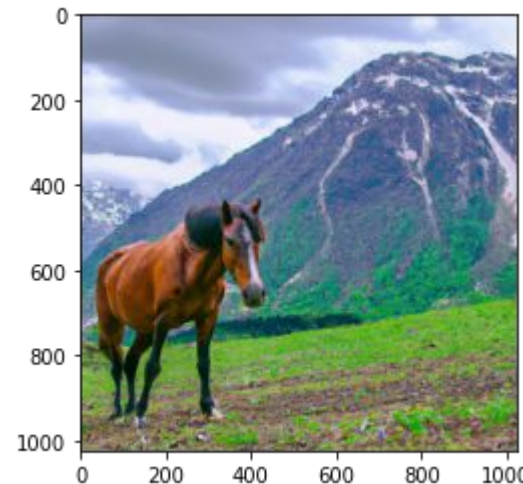
```
[0.1266768 0.12865311 0.1397564 0.1431053 0.15863107 0.14128713 0.16189025]
```

```
Best Prompt: mountain horse mountains field
```

```
Similarity: 0.87423253
```



Boston University



Results: *Prompt Generation with Target Image Only* (Policy Gradient Method)

```
# Hyperparameters
```

```
n_captions = 7, n_tokens = 8, max_episodes = 10, n_iterations = 20
```


```
prompt_length = 4
```

```
Prompt Space: ['of', 'is', 'sitting', 'sofa', 'couch', 'grey', 'long', 'cat']
```

```
Prompt Space Probabilities: [0.11573587 0.11826485 0.1274537 0.13060148 0.13342977  
0.11866169 0.12265412 0.13319854]
```

```
Best prompt:
```

```
Similarity:
```



Boston University

Limitations of ImagePrompt

- Relies on the performance and computation speed of other pre-trained models.
- Server on which DALL-E runs is unstable at times and crashes during training.
- DALL-E image generation is slow so training is slow.
- The model would have to be trained for every brand new image.

Future Works

- Train the model for more iterations and evaluate performance.
- Testing larger prompt sizes with more complex images.

Thank you!

Questions?